

Event Driven Architecture in TD Retail Platform




Sattria Sowanit,
Chief Software Architect,
TD Tawandang



Agenda

- TD Retail Business Journey
- Microservices & Event Driven Architecture
- Our Techniques
- Demo





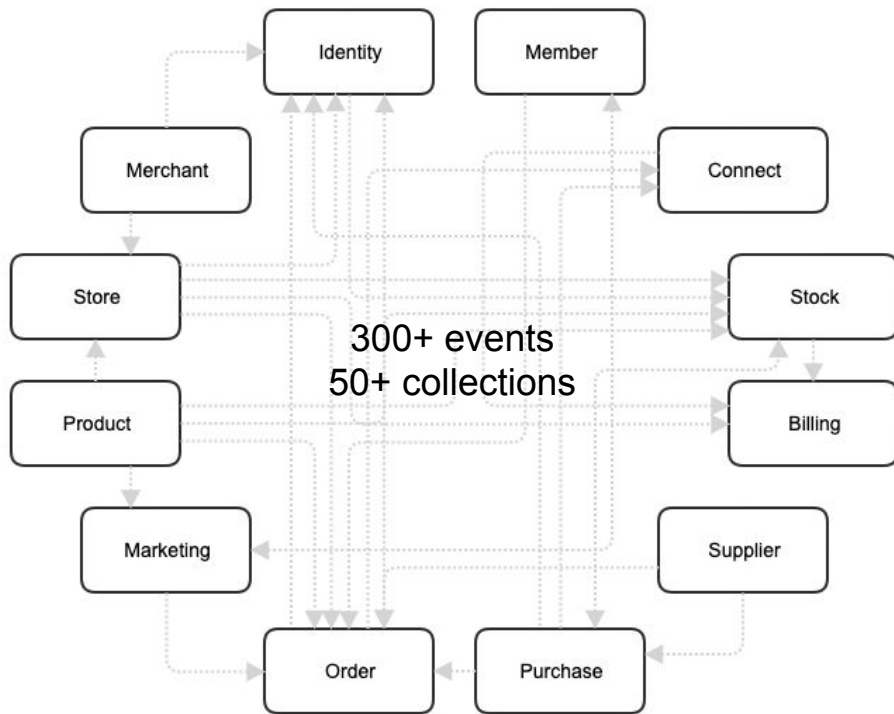
TD Retail Business Journey





Microservices & Event Driven Architecture

● Microservices



Clients

- 5K+ POS devices
- 5.5K+ handheld devices

Workloads

- 5K+ stores / 8 DC
- 800K+ bills / day
- 2K+ order / day
- 3.5M+ stock moving / day

Microservices

- 12 core modules
- 60+ sub-modules

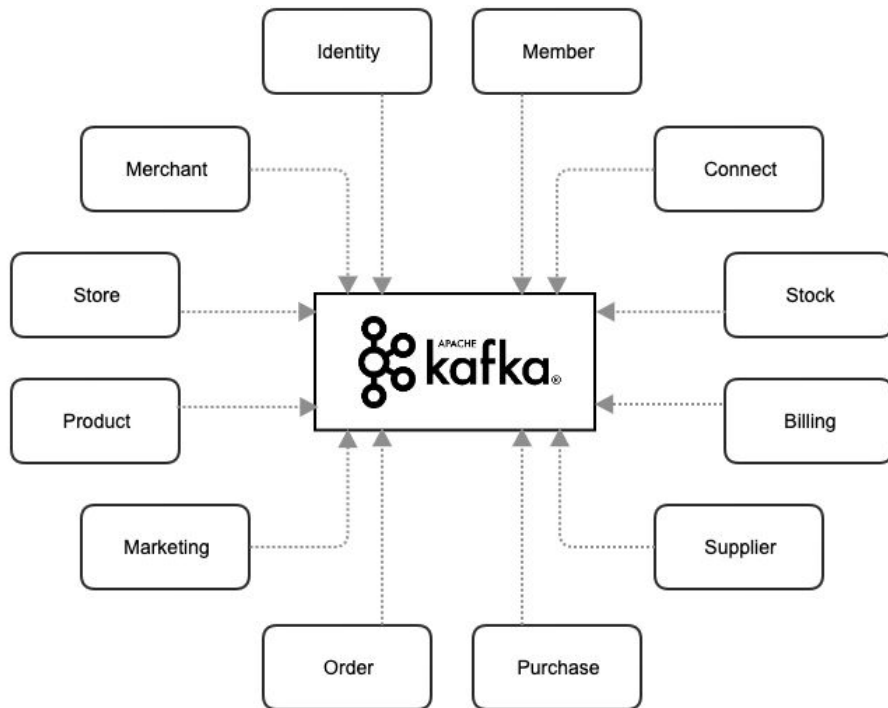
Runtime Platform

- Google Kubernetes Engine
- 30+ Worker Nodes
- 140+ Deployments
- 300+ Pods



Microservices & Event Driven Architecture

- **Event and Data Streaming**



Event Source

- 300+ event types
- 200+ records / second
- 30+ kilobytes / second
- 20+ connectors

Change Data Capture

- 50+ collections
- 200+ records / second
- 400+ kilobytes / second
- 30+ connectors


Runtime Platform

- Confluent Cloud
- Self-managed Connect Cluster
 - 20+ workers



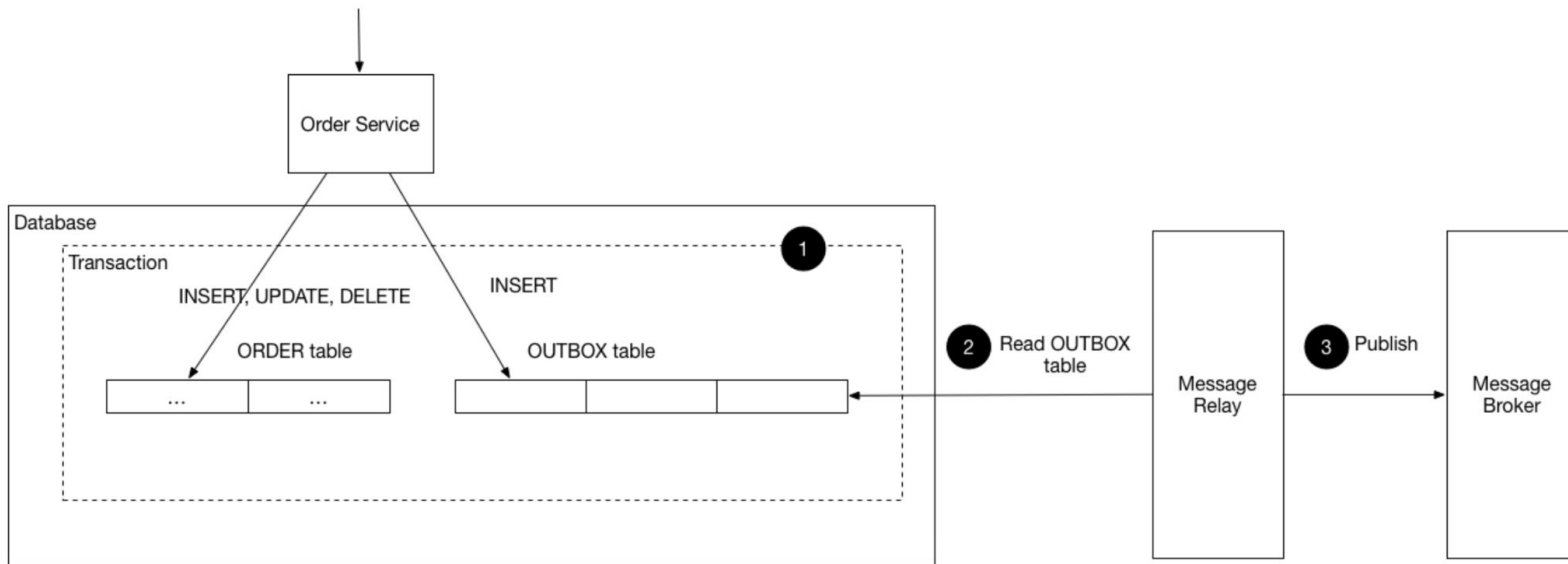
Our Techniques

- Outbox Pattern
- Debezium
- Kafka Connect
- Kafka Producer
- Configuration



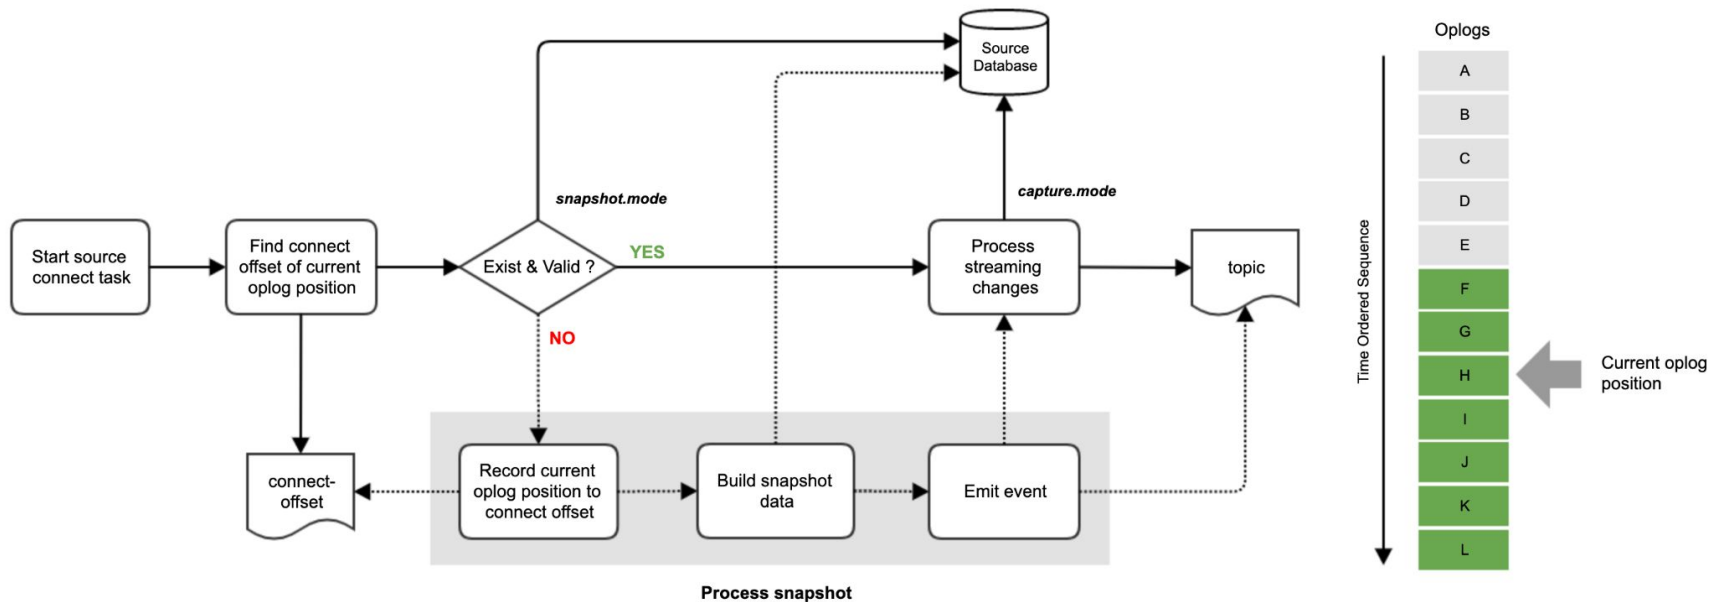


- **Data and Event Consistency**





- Source and Event Consistency



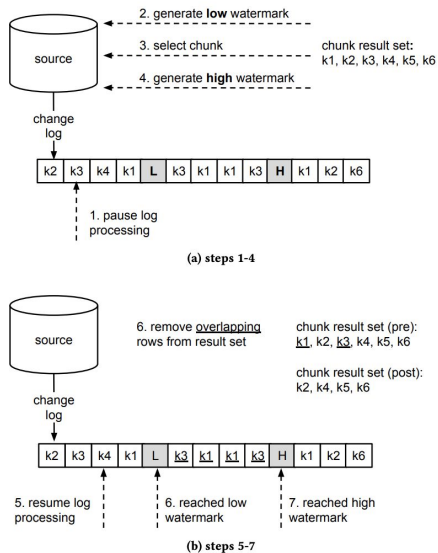


Figure 3: Watermark-based Chunk Selection




Figure 4: Order of output writes. Interleaving log capture with full data capture.

● Legacy Snapshot (Debezium <= 1.5)

- The near-impossibility of adding of additional tables to the captured tables list, if existing data must be streamed
- A **long-running process** for consistent snapshotting that cannot be terminated or resumed
- Change data streaming being **blocked** till the snapshot is completed

● Incremental Snapshot (Debezium >= 1.6)

- Watermark-based Snapshot (Netflix)
- Signalling Table

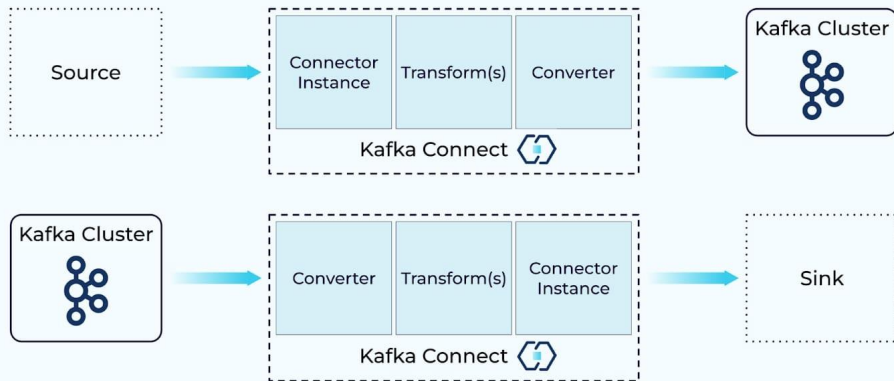
<https://arxiv.org/pdf/2010.12597v1.pdf>



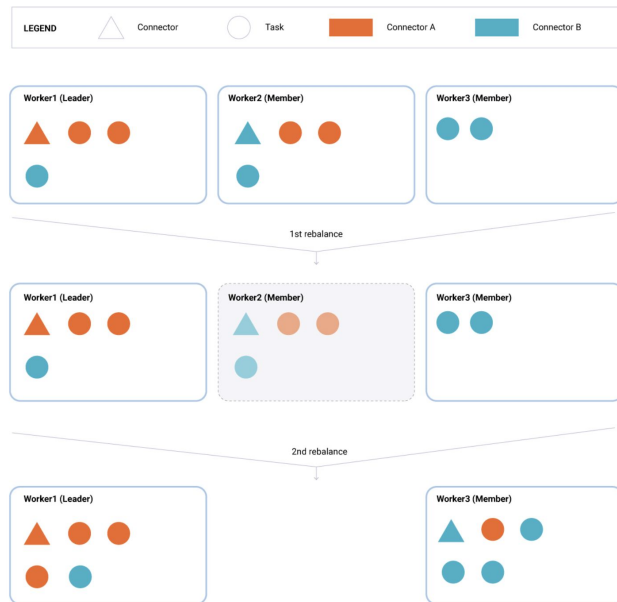
Techniques - Kafka Connect

- Scalability and Fault Tolerance

Inside Kafka Connect



@TheDanicaFine | developer.confluent.io




<https://developer.confluent.io/learn-kafka/kafka-connect/how-connectors-work/>



Techniques - Kafka Connect

- Stop the World Rebalancing (Kafka \leq 2.2)





● KIP-415 Incremental Cooperative Rebalancing (Kafka >= 2.3)



Goals

- Address the challenges at **large scale**

Why **incremental**?

- No need to reach final state within **single rebalance** round
- A **grace period** is configurable
- The protocol coverages smoothly to a state of balance load

Why **cooperative**?

- Resource **revocation** and **release** is graceful







Techniques - Kafka Producer

- Partition Strategy

Direct, Key Hashing, Round Robin (Kafka <=2.3) and Sticky (Kafka >=2.4)



p99 Latency: 3 Producers, 10,000 msg/sec (No Flush)





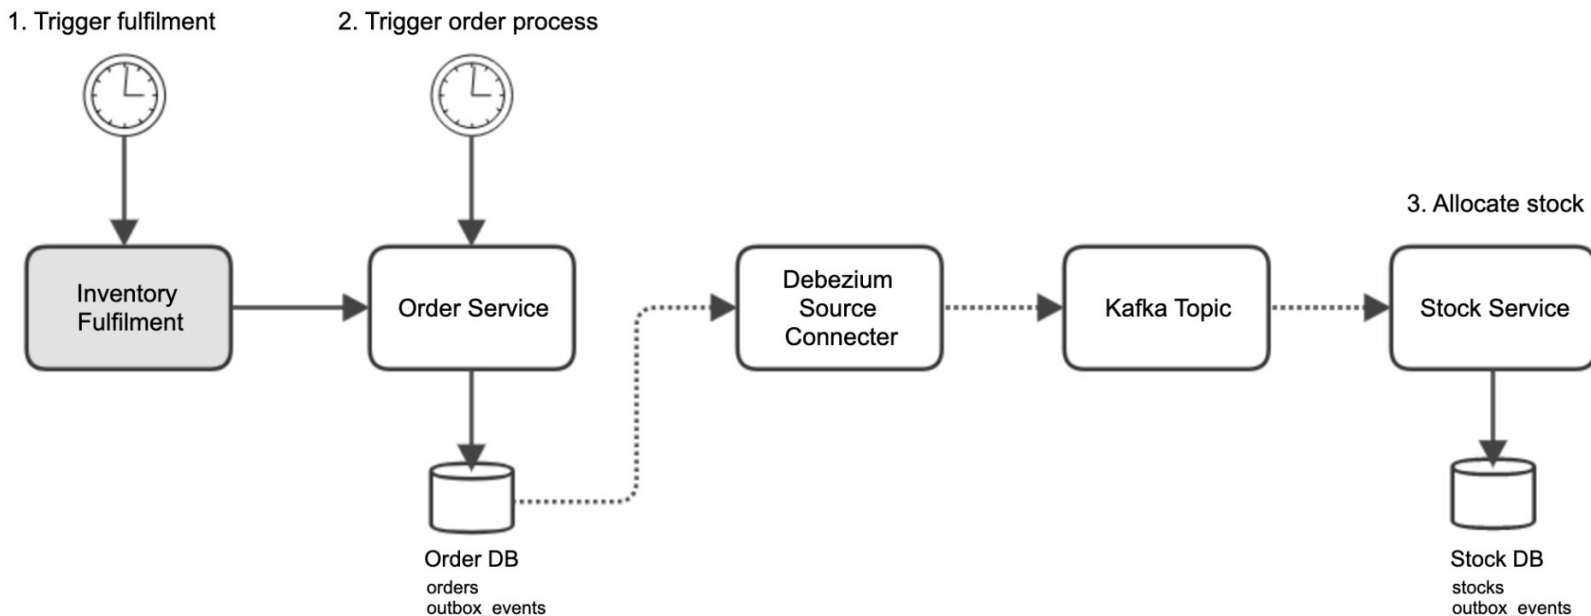
Techniques - Configuration

Type	Name	Used By	Default Value	Our Value
Connect	task.shutdown.graceful.timeout.ms	Fault tolerance	5000	15000
Connect	scheduled.rebalance.max.delay.ms	Incremental cooperation rebalance	300000	300000
Connect	offset.flush.timeout.ms	Large snapshot processing	5000	15000
Connect	offset.flush.interval.ms	Large snapshot processing	60000	30000
Producer	batch.size	Sticky partitioning	16384	524288
Producer	linger.ms	Sticky partitioning	0	5000
Producer	max.request.size	Large snapshot processing	1048576	8388608
Consumer	heartbeat.interval.ms	Fault tolerance	3000	5000
Consumer	session.timeout.ms	Fault tolerance	45000	60000
Consumer	max.poll.interval.ms	Fault tolerance	300000	600000
Consumer	max.poll.records	Fault tolerance	500	200



Demo - Use Case

- **Stock is allocated after order has been processed**





Demo - Event Source

2 usages · sattraso

```
public Mono<Order> saveOrder(Order order) {
    try {
        final OutboxEvent event = OutboxEvent.builder()
            .eventId(UUID.randomUUID().toString())
            .eventType(order.getStatus().getEventType())
            .aggregateType(OUTBOX_AGGREGATE_TYPE)
            .payload(mapper.writeValueAsString(order))
            .build();

        log.info("fire outbox_event: {}", event);
        if (order.getId() != null) {
            return template.inTransaction().execute(action ->
                action.save(order).zipWith(action.insert(event))) Flux<Tuple2<Order, OutboxEvent>>
                .map(tuple -> tuple.getT1()) Flux<Order>
                .next();
        }
        return template.inTransaction().execute(action ->
            action.insert(order).zipWith(action.insert(event))) Flux<Tuple2<Order, OutboxEvent>>
            .map(tuple -> tuple.getT1()) Flux<Order>
            .next();
    } catch (JsonProcessingException e) {
        return Mono.error(e);
    }
}
```




Demo - Connector

```
{
  "name": "order-outbox-source",
  "config": {
    "connector.class": "io.debezium.connector.mongodb.MongoDbConnector",
    "mongodb.name": "mongodb.local",
    "mongodb.hosts": "mongodb://mongo:30001,mongo2:30002,mongo3:30002/order?replicaSet=rs0",
    "transforms": "router",
    "transforms.router.type": "com.example.outbox.Router",
    "transforms.router.topic": "outbox_events",
    "database.whitelist": "order",
    "collection.whitelist": "order[.]outbox_events",
    "tasks.max": "1"
  }
}
```



Demo - Consumer

```
± satraso
@PostConstruct
public void subscribe() {
    receiver.receive()
        .delayUntil(record -> handle(record))
        .doOnSuccess(o -> record.receiverOffset().acknowledge())
        .doOnError(e -> log.error("[{}] error occurred => {}", record.key(), e.getMessage()))
    }
    .subscribe();
}
```

```
± satraso
@PreDestroy
public void destroy() { log.warn("Shutting down stream ..."); }
```

1 usage ± satraso

```
private Mono<Void> handle(ReceiverRecord<String, Event> record) {
    log.info("Received message from topic: {}, key:{}, value:{}", record.topic(), record.key(), record.value());
    String eventType = record.value().getEventType().toString().toUpperCase();
    switch (EventType.valueOf(eventType)) {
        case ORDER_PROCESSING:
            log.info("Handle event_type: {}", eventType);
            String payload = record.value().getPayload().toString();
            return Mono.fromCallable(() -> mapper.readValue(payload, Order.class)) Mono<Order>
                .flatMap(o -> service.allocateStock(o)) Mono<Void>
                .then();
        default:
            log.info("Un-handle event_type: {}", eventType);
    }
    return Mono.empty();
}
```



<https://github.com/Tookdee88/spring-boot-outbox-pattern>